

NeuroCortex

Self-Reinforcing Development Framework (SRDF)

A Conceptual White Paper

Author: Mohammed Qaid Al-Athwary
ORCID: 0009-0006-9075-072X
Email: mohammedqaidalathwary@gmail.com

Version: 1.0
Date: October 2026
Document type: Conceptual White Paper

1. Abstract

This white paper introduces **NeuroCortex**, a research direction centered on the **Self-Reinforcing Development Framework (SRDF)**. SRDF proposes a controlled, auditable approach to runtime structural adaptation in machine learning systems. Rather than allowing unbounded self-modification, SRDF separates the adaptation process into explicit stages: observation, analysis, candidate generation, evaluation, authorization, and controlled state transition.

This document describes the conceptual foundations of the framework, its principal components, and its design principles. It is a conceptual document only. It does not report experimental results, does not present quantitative evaluation, and does not make comparative performance claims.

2. Motivation

Contemporary machine learning systems operate under a fixed architecture once deployed. When the data distribution changes - a phenomenon widely observed in production environments - the model typically loses effectiveness over time. Recovering performance generally requires human intervention: retraining, re-engineering features, or redesigning the pipeline.

Recent automation efforts, such as automated machine learning (AutoML) pipelines, reduce some of this burden by searching within predefined spaces of models and hyperparameters. However, these approaches operate within fixed architectural boundaries. They can tune what is given, but they cannot propose new structural forms.

The central question this project explores is conceptual: **can a system propose and adopt structural modifications to its own learning pipeline at runtime, in a controlled and auditable way?** SRDF is a proposed conceptual framework for approaching this question. It is not a claim about capability, and it is not a claim about superiority over existing methods.

3. The SRDF Framework

SRDF is organized around three principal components, each with a distinct role.

3.1 Trawler

The Trawler observes the current runtime context and identifies conditions that may justify adaptation. It analyzes available signals and reports findings to the framework. Its role is observation and analysis, not decision-making.

3.2 Generator

The Generator produces candidate structural or learning modifications. Candidates are proposals, not commitments. Each candidate describes a possible change to the current pipeline or learning configuration. The Generator does not evaluate its own candidates and does not authorize them.

3.3 Arbiter

The Arbiter evaluates candidates against predefined constraints and authorizes only those that satisfy the required conditions. Its role is independent of the Generator. A candidate that does not pass the Arbiter is rejected and does not affect the runtime state.

3.4 Workflow

The workflow is ordered and stage-separated:

```
OBSERVE -> ANALYZE -> GENERATE -> EVALUATE -> AUTHORIZE  
-> COMMIT_OR_REJECT -> RECORD
```

Each stage has a defined role and a defined output. Stage separation is a design principle, not a runtime detail.

4. Design Principles

The framework is guided by the following principles.

4.1 Separation of generation and authorization

The component that proposes a change is never the component that approves it. This separation is intended to preserve the integrity of the decision process.

4.2 Bounded structural change

Structural modification is permitted only through predefined candidate forms. Unrestricted program synthesis or unconstrained architectural generation is outside the scope of the framework.

4.3 Explicit authorization

Every structural change requires an explicit authorization step. A candidate that is not authorized cannot alter the runtime state.

4.4 Auditable records

The framework records the observed context, the candidates generated, their evaluation, their authorization outcome, and the resulting state transition. Records are intended to be traceable.

4.5 Protocol before execution

When the framework is applied to a specific experimental setting, the protocol is defined in advance. Modifications to the protocol after execution are not permitted; any methodological change requires a new protocol version.

5. Scope and Boundaries

This document is **conceptual**. Its purpose is to describe a research direction and its design principles.

This document does **not** claim:

- general or universal performance improvement,
- statistical significance of any result,
- generalization beyond a specific experimental setting,
- self-evolving artificial general intelligence,

- unrestricted autonomy or unbounded self-modification,
- superiority over any existing method or architecture.

This document is **not a research manuscript**. It is not a preprint of a future manuscript, and it does not contain experimental results. Any future empirical claims will be reported separately in a peer-reviewed publication.

6. Project Status

The NeuroCortex project is **under active development**. The current repository hosts:

- the conceptual framework described in this document,
- the protocols and design principles associated with it,
- the organizational structure of the experimental phases,
- artifacts related to the phases that have been executed.

The repository is intended to preserve the framework, its protocols, and its evidence structure in a form that is traceable and reproducible. Descriptions of completed phases are maintained in their respective directories. This white paper does not summarize those phases and does not report their results.

7. Ethics and Governance

The prospect of runtime structural adaptation raises questions that are both technical and ethical. Two points are stated explicitly here.

First, any adaptation mechanism must operate under transparent rules. The framework therefore requires that observation, generation, evaluation, and authorization be separated and recorded.

Second, any adaptation mechanism must operate within explicit boundaries. The framework does not permit unbounded modification, and it does not claim to be autonomous in any general sense. The scope of adaptation is bounded by the protocol under which the framework is executed.

8. Contact

Author: Mohammed Qaid Al-Athwary

ORCID: 0009-0006-9075-072X

Email: mohammedqaidalathwary@gmail.com

Repository: github.com/mohammedqaidamanlift-hub/-NEUROCORTEX

9. Citation

If you reference this white paper, please cite it as:

Al-Athwary, M. Q. (2026). NeuroCortex: Self-Reinforcing Development Framework (SRDF) - A Conceptual White Paper (Version 1.0). GitHub.
<https://github.com/mohammedqaidamanlift-hub/-NEUROCORTEX>

A persistent identifier (DOI) will be assigned through Zenodo upon archival.

End of white paper.